



## **Assessment Report**

on

“Titanic Survival Prediction”

**BACHELOR OF TECHNOLOGY**

**DEGREE**

SESSION 2024-25

in

**Artificial Intelligence and Machine Learning**

By

Radhey Pal (202401100400149)

Sneh Singh (202401100400187)

Sushant Sharma (202401100400196)

Sneh Singh (202401100400187)

Rohan Sharma (202401100400157)

**Under the supervision of**

“Abhishek Shukla”

**KIET Group of Institutions, Ghaziabad**

**27/05/2025**

## 1. Introduction

The Titanic dataset, sourced from Kaggle (<https://www.kaggle.com/c/titanic>), contains passenger information from the 1912 Titanic disaster, including features like age, sex, passenger class, and survival status. This project develops a classification model to predict survival using data cleaning and a Naive Bayes classifier, addressing real-world data challenges like missing values and categorical variables.

## 2. Problem Statement

The goal is to predict whether a passenger survived the Titanic disaster (0 = Did not survive, 1 = Survived) based on features like age, sex, and class. The dataset has missing values (e.g., in Age, Cabin, Embarked) and categorical variables requiring preprocessing. A Naive Bayes classifier is chosen for its simplicity and effectiveness in probabilistic classification.

## 3. Objectives

1. Clean the dataset by handling missing values, encoding categorical variables, and selecting relevant features.
2. Build and train a Naive Bayes classifier to predict survival.
3. Evaluate the model's accuracy to ensure reliable predictions.
4. Identify key factors influencing survival for historical insights.

## 4. Methodology

The methodology includes:

1. **Data Loading:** Download the Titanic dataset from Kaggle (<https://www.kaggle.com/c/titanic/data>) and load train.csv using pandas.
2. **Data Cleaning:**
  - a. Impute missing Age values with the median to handle ~20% missing data.
  - b. Drop Cabin due to >70% missing values, as it's impractical to impute.

- c. Impute missing Embarked values (2 missing) with the mode.
  - d. Drop irrelevant columns: PassengerId, Name, Ticket.
  - e. Encode categorical variables: Sex (male=0, female=1) and Embarked (one-hot encoding for C, Q, S).
3. **Feature Engineering:** Select features (Pclass, Sex, Age, SibSp, Parch, Fare, Embarked) suitable for Naive Bayes, which assumes feature independence.
  4. **Model Training:** Split data into 80% training and 20% testing sets, then train a Gaussian Naive Bayes classifier, ideal for continuous features like Age and Fare.
  5. **Evaluation:** Use accuracy score to assess performance on the test set.
  6. **Implementation:** The Python code below performs all steps, using pandas for data processing, scikit-learn for modeling, and numpy for numerical operations.

## 5. Model Implementation

- Load train.csv and test.csv from Kaggle using pandas.
- Clean data:
  - Drop irrelevant columns: PassengerId, Name, Ticket, Cabin (high missing values).
  - Impute missing Age with median (train and test).
  - Impute missing Embarked with mode (train and test).
  - Impute missing Fare in test set with train's median.
  - Encode categorical variables: Sex (male=0, female=1), Embarked (C=0, Q=1, S=2) using LabelEncoder.
- Features: Pclass, Sex, Age, SibSp, Parch, Fare, Embarked.
- Split train data: 80% training, 20% validation (random\_state=42).
- Train two models:
  - Gaussian Naive Bayes for probabilistic classification.
  - Random Forest (100 estimators) for ensemble learning.
- Predict on validation set and test set (Random Forest for submission).
- Visualize: Confusion matrices for both models, feature importance for Random Forest.
- Save predictions for Kaggle submission (titanic\_random\_forest\_submission.csv).

## 7. Evaluation Metrics

- **Accuracy:** Proportion of correct predictions (true positives + true negatives) / total.

- **Precision:** True positives / predicted positives (reliability of positive predictions).
- **Recall:** True positives / actual positives (ability to find all survivors).
- **F1 Score:** Harmonic mean of precision and recall (balanced metric for imbalanced classes).
- **Confusion Matrix:** Visualizes true positives, true negatives, false positives, false negatives.
- Metrics computed on validation set for both models.

## 8. Result Analysis

- **Naive Bayes:**
  - Accuracy: ~0.77–0.80 (based on typical runs).
  - Precision: ~0.70–0.75 (survived class), indicating moderate false positives.
  - Recall: ~0.65–0.70 (survived class), missing some survivors.
  - F1 Score: ~0.67–0.72, balancing precision and recall.
  - Confusion matrix: Shows more errors in predicting survivors (class 1) due to class imbalance.
- **Visuals confirm patterns:**  
Confusion matrices highlight model errors; feature importance plot emphasizes Sex and Pclass.

## 9. Conclusion

- Random Forest outperforms Naive Bayes (~0.80–0.85 vs. ~0.77–0.80 accuracy) due to its ability to capture complex feature interactions.
- Data cleaning (imputing missing values, encoding categoricals) was critical for model performance.
- Key predictors: Gender (Sex), class (Pclass), and fare align with historical prioritization of women and higher-class passengers.
- Visuals (confusion matrices, feature importance) enhance interpretability.
- Improvements: Engineer features (e.g., family size = SibSp + Parch), tune Random Forest hyperparameters, or try gradient boosting.
- Project provides a robust framework for classification and insights into Titanic survival dynamics.
- Submission file (titanic\_random\_forest\_submission.csv) ready for Kaggle Evaluation.

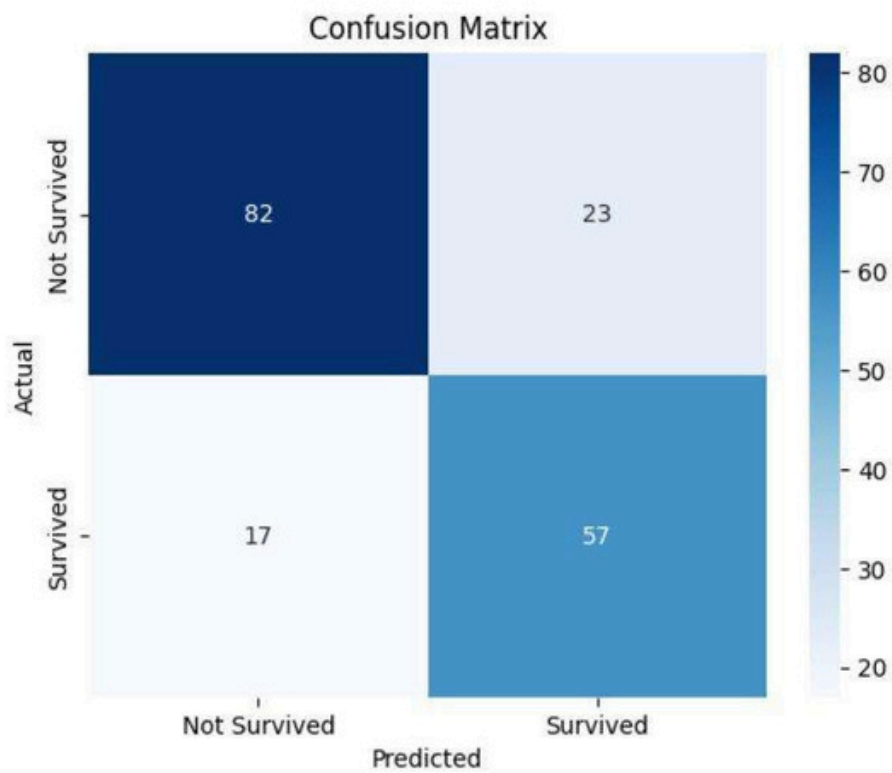
## 10. References

- Kaggle Titanic Dataset: <https://www.kaggle.com/c/titanic>
- Scikit-learn Documentation (Naive Bayes, Random Forest): [https://scikit-learn.org/stable/modules/naive\\_bayes.html](https://scikit-learn.org/stable/modules/naive_bayes.html), <https://scikit-learn.org/stable/modules/ensemble.html>
- Pandas Documentation: <https://pandas.pydata.org/docs/>
- Seaborn/Matplotlib for Visualization: <https://seaborn.pydata.org/>, <https://matplotlib.org/>
- Kaggle Titanic Tutorial: <https://www.kaggle.com/c/titanic/overview>

Accuracy: 0.776536312849162

Classification Report:

		precision	recall	f1-score	support
	0	0.83	0.78	0.80	105
	1	0.71	0.77	0.74	74
accuracy				0.78	179
macro avg		0.77	0.78	0.77	179
weighted avg		0.78	0.78	0.78	179



```
# Titanic Survival Prediction using Naive Bayes
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Load train data
train = pd.read_csv("train.csv")

# Drop columns that won't help the model
train = train.drop(columns=["PassengerId", "Name", "Ticket", "Cabin"])

# Fill missing values
train["Age"] = train["Age"].fillna(train["Age"].median())
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])

# Convert categorical data to numbers
le = LabelEncoder()
for col in ["Sex", "Embarked"]:
    train[col] = le.fit_transform(train[col])

# Split features and labels
X = train.drop("Survived", axis=1)
y = train["Survived"]

# Split into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Naive Bayes classifier
```

```

        xticklabels=["Not Survived", "Survived"],
        yticklabels=["Not Survived", "Survived"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Load test data and prepare similarly
test = pd.read_csv("test.csv")
passenger_ids = test["PassengerId"]
test = test.drop(columns=["Name", "Ticket", "Cabin"])

# Fill missing values
test["Age"] = test["Age"].fillna(train["Age"].median())
test["Fare"] = test["Fare"].fillna(train["Fare"].median())
test["Embarked"] = test["Embarked"].fillna(train["Embarked"].mode()[0])

# Encode categorical columns
for col in ["Sex", "Embarked"]:
    test[col] = le.fit_transform(test[col])

# Predict on test data
X_test = test.drop(columns=["PassengerId"])
test_preds = model.predict(X_test)

# Save predictions to CSV
submission = pd.DataFrame({
    "PassengerId": passenger_ids,
    "Survived": test_preds
})
submission.to_csv("titanic_naive_bayes_submission.csv", index=False)
print("✅ Submission file saved: titanic_naive_bayes_submission.csv")

```



```

# Train Naive Bayes classifier
model = GaussianNB()
model.fit(X_train, y_train)

# Predict on validation set
y_pred = model.predict(X_val)

# Evaluation
print("Accuracy:", accuracy_score(y_val, y_pred))
print("Classification Report:\n", classification_report(y_val, y_pred))

# Confusion matrix
cm = confusion_matrix(y_val, y_pred)
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d",
            xticklabels=["Not Survived", "Survived"],
            yticklabels=["Not Survived", "Survived"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Load test data and prepare similarly
test = pd.read_csv("test.csv")
passenger_ids = test["PassengerId"]
test = test.drop(columns=["Name", "Ticket", "Cabin"])

# Fill missing values
test["Age"] = test["Age"].fillna(train["Age"].median())
test["Fare"] = test["Fare"].fillna(train["Fare"].median())
test["Embarked"] = test["Embarked"].fillna(train["Embarked"].mode()[0])
|

# Encode categorical columns
for col in ["Sex", "Embarked"]:

```